

ML - 9 Logistic Regression (24 Aug 22)

Wednesday, August 10, 2022 1:23 PM

Pre-Class

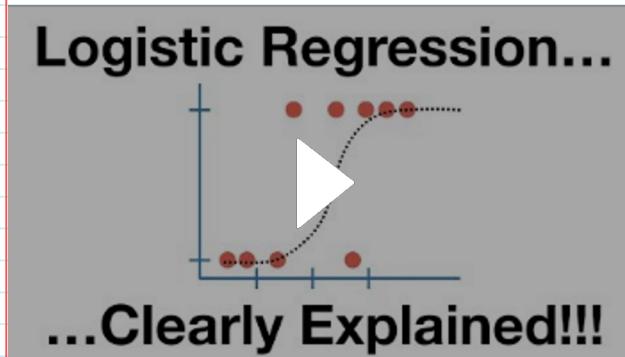
Logistic Regression

Introduction

The linear regression fits a line and is used to predict a continuous value. For example, we estimate the price of a house using linear regression. However, the logistic regression fits a **s-shaped** (sigmoid) function and predicts whether something is true or false. It is mainly used in **binary classification problems**. In the video below, you will see the logic behind logistic regression and examine a binary classification problem related to mice being obese or not obese.

★ Watch:

[StatQuest: Logistic Regression](#)



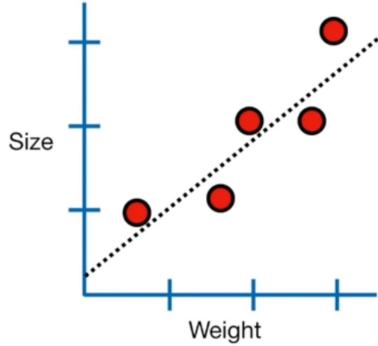
Interview Q&A



Q: Why can't **linear regression** be used in place of logistic regression for binary classification?

A: In **linear regression**, the output is continuous. In case of binary classification, an output of a continuous value does not make sense. For binary classification problems, **linear regression** may predict values that can go beyond 0 and 1. If we want the output in the form of probabilities, which can be mapped to two different classes, then its range should be restricted to 0 and 1. As the logistic regression model can output probabilities with logistic/sigmoid function, it is preferred over **linear regression**.

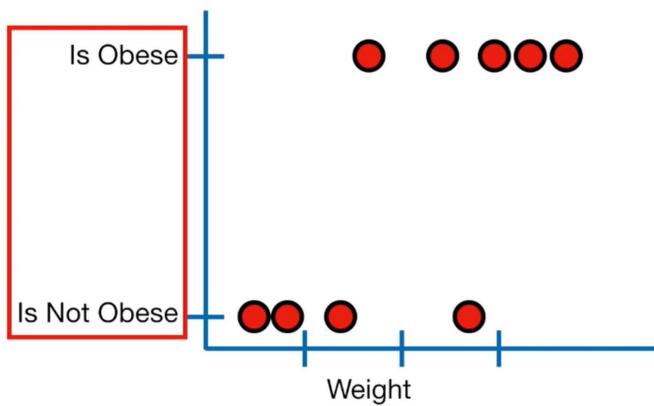
Linear Regression:



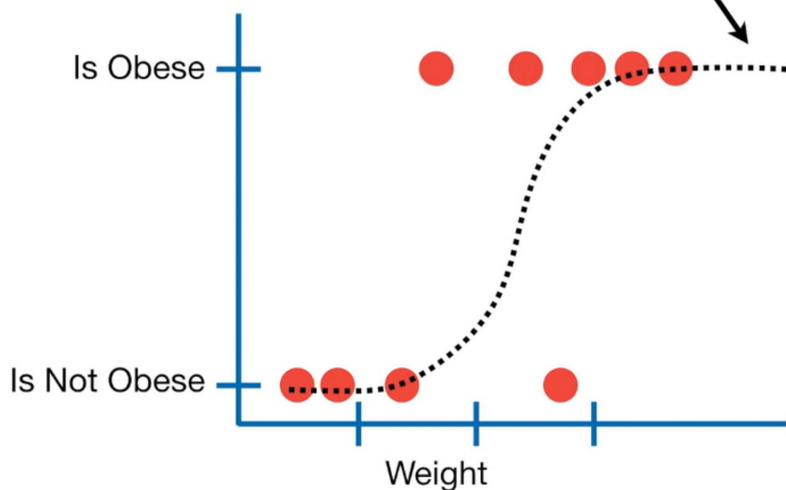
...and with that line, we could do a lot of things:

- 1) Calculate R^2 and determine if **weight** and **size** are correlated. Large values imply a large effect.
- 2) Calculate a p-value to determine if the R^2 value is statistically significant.
- 3) Use the line to predict **size** given **weight**.

Logistic regression predicts whether something is **True** or **False**, instead of predicting something continuous like **size**.

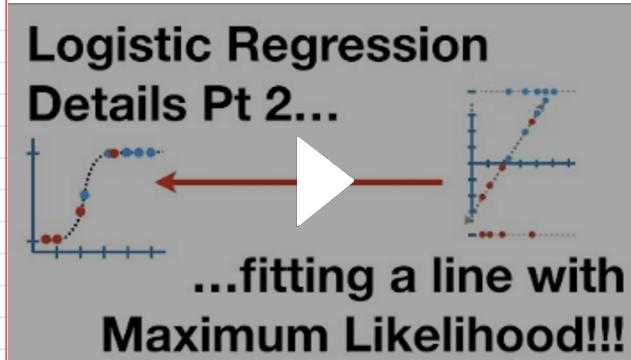


...also, instead of fitting a line to the data, logistic regression fits an "S" shaped "logistic function".



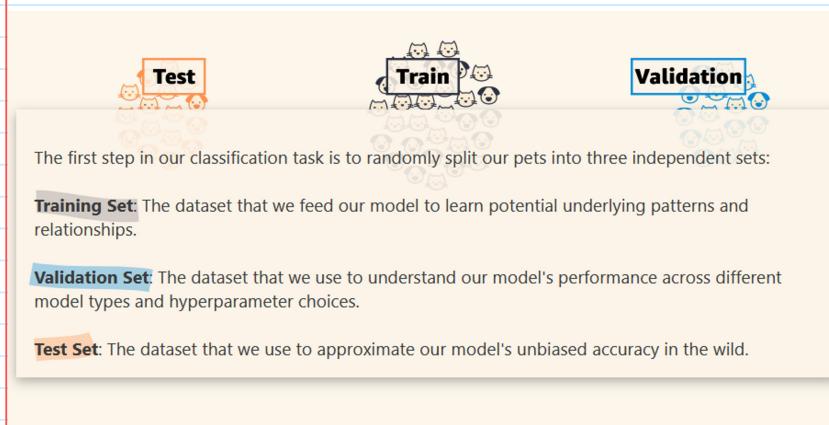
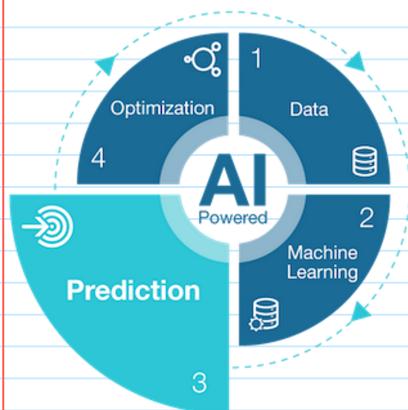


Watch:
[Logistic Regression Details Pt 2: Maximum Likelihood](#)



Evaluating Performance-Classification Error Metrics

In a binary classification model, our prediction can only get two results: **correct** or **incorrect**. We use classification error metrics to measure and evaluate the performance of our predictions.



The major classification metrics:

- **Accuracy:** The proportion of correct predictions (both true positives and true negatives) among the total number of cases examined.
- **Recall:** Also known as sensitivity. It is the fraction of relevant instances that were retrieved.
- **Precision:** Also called positive predictive value. It is the fraction of relevant instances among the retrieved instances
- **F1-Score:** The F1 score is the harmonic mean of the precision and recall.



Watch:
[Never Forget Again! // Precision vs Recall with a Clear Example of Precision and Recall](#)



Confusion Matrix:

The classification metrics mentioned above (accuracy, precision, etc.) are generated from the confusion matrix. A confusion matrix is a frequently used table to describe the performance of a classification model on a series of test data where real values are known.

💡 Tips:

- Keep in mind that understanding the confusion matrix is critical in the performance measurement of classification models.



Watch:
[The NO CONFUSION matrix! // What is the confusion matrix? // Confusion matrix visual explanation](#)



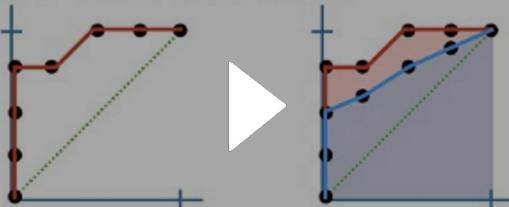
Receiver operating characteristic (ROC) curve:

This is a graphical plot that illustrates the performance of a binary classifier as its discriminant threshold is varied. The curve is created by plotting the true positive rate (TPR) against the false positive rate (FPR) at various threshold values.



Watch:
[ROC and AUC, Clearly Explained!](#)

ROC and AUC....



...Clearly Explained!!!



Interview Q&A

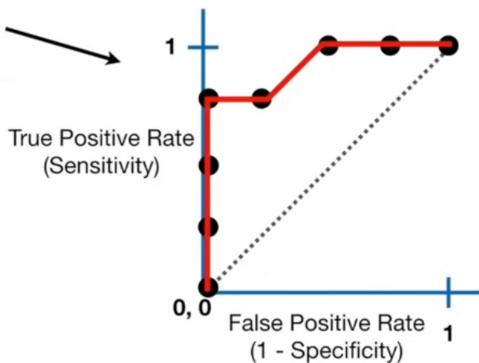


Q: Why is accuracy not a good measure for classification problems?

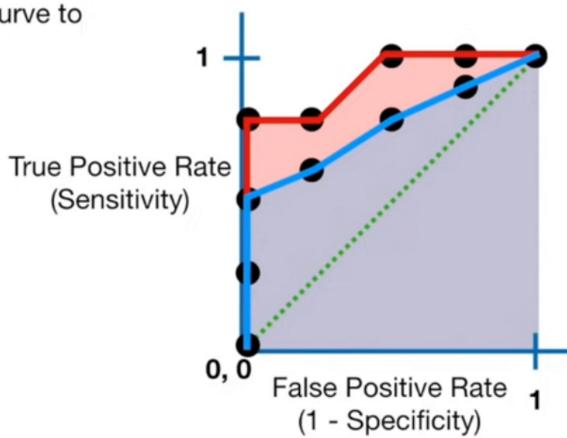
A: Accuracy is not a good measure for classification problems because it gives equal importance to both false positives and false negatives.

However, this may not be the case in most business problems. For example, in case of cancer prediction, declaring cancer as benign is more serious than wrongly informing the patient that he is suffering from cancer. Accuracy gives equal importance to both cases and cannot differentiate between them.

So instead of being overwhelmed with confusion matrices, **Receiver Operator Characteristic (ROC)** graphs provide a simple way to summarize all of the information.

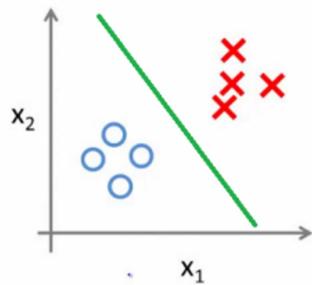


The **AUC** makes it easy to compare one **ROC** curve to another.

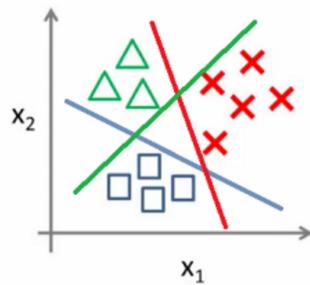


Logistic Regression - Multi-class Classification with Python

Binary classification:

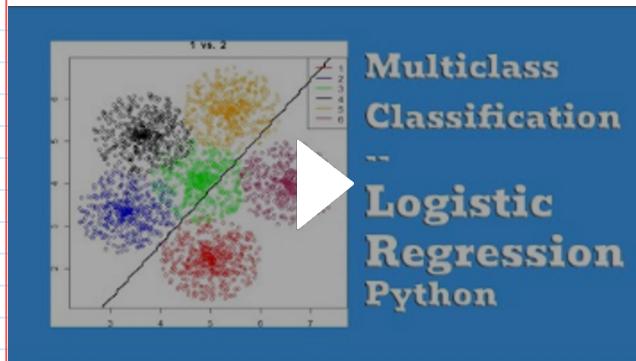


Multi-class classification:

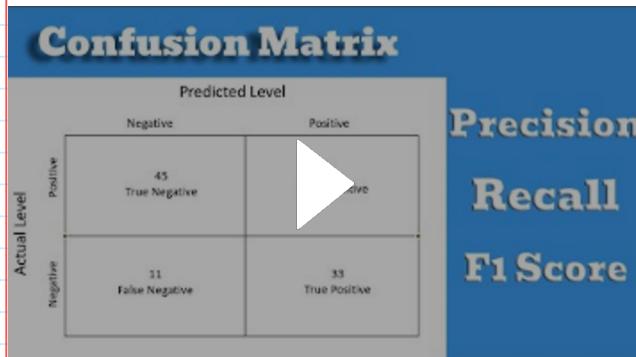


★ Watch:

[Multiclass Classification With Logistic regression in Python | Sklearn](#)



[Precision, Recall and F1 Score for Multiclass Classification - sklearn | Python](#)



★ Watch:

💡 Interview Q&A

In-Class

Stratify = aynı dağılımda bölmek

Official documentation:

https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html?highlight=train_test_split

stratify array-like, default=None

If not None, data is split in a stratified fashion, using this as the class labels. Read more in the User Guide.

Some classification problems can exhibit a large imbalance in the distribution of the target classes: for instance there could be several times more negative samples than positive samples. In such cases it is recommended to use stratified sampling as implemented in StratifiedKFold and StratifiedShuffleSplit to ensure that relative class frequencies are approximately preserved in each train and validation fold.

This **stratify** parameter makes a split so that the proportion of values in the sample produced will be the same as the proportion of values provided to parameter **stratify**.

For example, if variable **y** is a binary categorical variable with values 0 and 1 and there are 25% of zeros and 75% of ones, **stratify=y** will make sure that your random split has 25% of 0's and 75% of 1's.

Michael Gd-Tech.Coord.:

stratify target değerlerinin iki tarafa aynı oranda dağıtılmmasını sağlıyor. yani target value(0,1) oranları mesela yüzde 30 (0) yüzde 70(1) olsun. train test split random olarak bölme yaptığı için birler bi tarafa sıfırlar bi tarafa toplanmasın diye stratify kullanılır. sonuçta bölme işlemi nihayetlenince train set yüzde 30 (0) yüzde 70(1) target value sahip olur aynı şekilde test set yüzde 30 (0) yüzde 70(1) target value sahip olur. target value unbalanced olduğunda kullanılır. diğer durumlarda da kullanılmasının sakıncası yoktur.

Soru:

burada neden pred_proba[:, 1] yaptıktı yani sadece 1. sütunu aldık? 0. sütunda de ne vardı pred_proba nın

Michael Gd-Tech.Coord.:

- model her zaman default 1 lerini hedef olarak görür ve hesaplamayı ona göre yapar. o yüzden 1. sütuna bakıyoruz
- Sınıfın 0 olma olasılığı

Soru:

0 ve 1 için metrikler çok farklı çıktı neden?

Orion - ML Instructor:

- 0 sınıfından fazla olduğundan ondan daha çok öğrenmiş
- imbalance data setlerinde modelin yanlışlığı fazla olan classa yönelir bu yüzdede scoreları genelde iyi çıkarlar.

Orion - ML Instructor:

parametre = modelin bulduğu
hyper parametre = bizim ayarladığınız

Limited-memory BFGS (**L-BFGS** or LM-BFGS) is an optimization algorithm in the family of quasi-Newton methods that approximates the Broyden–Fletcher–Goldfarb–Shanno algorithm (BFGS) using a limited amount of computer memory. It is a popular algorithm for parameter estimation in machine learning.

Orion - ML Instructor:

bu şekilde yazmak
`fp_rate, tp_rate, thresholds = roc_curve(y_train, y_pred_proba[:, 1])`
işte böyle.

John - ML Instructor:

modele verdigimiz yeni verinin feature sırası modeldeki ile aynı olmalı ve scaled olmalı

Only Label / target may have categorical value. Not true for features, be careful!

ROC is used in binary classification scenarios, not in multi-class.

ROC, AUC, precision, and recall visually explained – <https://paulvanderlaken.com/2019/08/16/roc-auc-precision-and-recall-visualy-explained/>

Orion - ML Instructor:

- modeliniz hangi sample range inde eğitilmiş ise o rangede en iyi çalışır
- lineerde alpha büyükçe regülarizasyon artırdı logistikte ise küçündükçe artıyor



In-Class Slides:

ML - S8/S9 - (Logistic Regression and Classification Error Metrics)

<https://app.peardeck.com/review/student/tbytsetdb>



In-Class Record:

ML - S9 - 24/08/22 John (Logistic Regression with Phyton)

<https://lms.clarusway.com/mod/page/view.php?id=24040>



Post-Class

Scikit-Learn: <https://scikit-learn.org/stable/index.html>

Yellowbrick: <https://www.scikit-yb.org/en/latest/quickstart.html>

Orion - ML Instructor:

Drop first

Although the category you choose to drop won't affect the model's performance, it can have a significant impact on the interpretability of the model.

If you plan to use the coefficients in your model to make accurate inferences about:

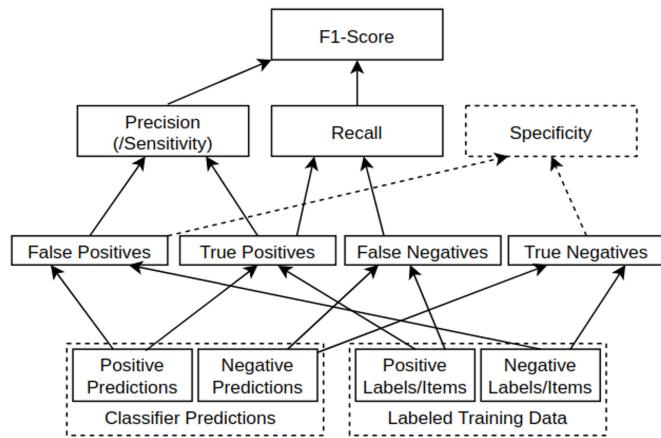
- impact on the target per predictor unit,
- or the magnitude of impact each predictor has in relation to others,

then you should carefully consider which categorical values you drop from your model.

A linear regression model's coefficients are interpreted in the context of a baseline model. For continuous variables, the baseline uses a reference point of 0. But for categorical variables, whichever column is dropped becomes the reference point, which has a significant impact on how coefficients are interpreted.

		Actual (True) Values	
		Positive	Negative
Predicted Values	Positive	TP	FP
	Negative	FN	TN

Confusion Matrix. Image by Author.



Hierarchy of Metrics from raw measurements / labeled data to F1-Score. Image by Author.

https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html



ML-9 Recap

M-G Logistic Regression

Phases

1) Data preparing / EDA

2) Divide data → features

→ label (target)

3) Train / Test split:

$X_{\text{train}}, X_{\text{test}}, y_{\text{train}}, y_{\text{test}} = \text{train_test_split}(X, y, \dots)$

Phases/steps build on top of each other.
The output of previous step is the input of next step.

$X = \text{df.drop}["\text{target"}], \text{axis}=1$

$y = \text{df}["\text{target"}]$

4) Scale X (we don't scale "y", since it's single column)

$X_{\text{train-scaled}} = \text{scaler}.\underline{\text{fit_transform}}(X_{\text{train}})$ - we fit-transform only X_{train}

$X_{\text{test-scaled}} = \text{scaler}.\underline{\text{transform}}(X_{\text{test}})$ - we transform X_{test} just

5) Select model (`LinearRegression()`, `LogisticRegression()` ...)

6) Train ("Fit") model with train data

`log-model.fit(X_train-scaled, y-train)`

7) Make predictions (by using $X_{\text{test-scaled}}$ and "trained" model)

$y_{\text{pred}} = \text{log-model.predict}(X_{\text{test-scaled}})$

8) Get probability estimates with "predict-proba"

$y_{\text{pred-proba}} = \text{log-model.predict_proba}(X_{\text{test-scaled}})$

9) Compare test data with prediction and probability.

`test-data = pd.concat([X-test, y-test])`

`test-data['pred'] = y-pred`

`test-data['pred-proba'] = y-pred-proba[:, 1]`

prob. estimates
of outcomes(1)
- if we put 0
we get prob. of
outcomes 0%

10) Get metrics and evaluate

- confusion matrix

- classification report

		True Neg	False Pos
Actual	True Neg	TN	FP
	False Pos	FN	True Pos
	Predict		

13) ROC and AUC review

14) Find best threshold

15) Final Model and deployment
train with all X data

11) Cross validate with train data
→ get average of scores.

12) Grid Search CV → to find best hyper-parameters

define parameters

`grid-model.fit(X-train-scaled, y-train)`

`grid-model.best_params_` best "C", "penalty", "solver", "class_weight"